

Exercise 4: Convolutional Neural Networks (CNNs)

Objective:

Apply the concepts of Convolutional Neural Networks (CNNs) as explained in the provided notebook to a new dataset. You will build, train, and evaluate a CNN model on one of the datasets listed below.

Instructions:

1. **Choose a Dataset:** Select one of the following datasets to work with:
 - a. **Fashion-MNIST:** A dataset of Zalando's article images—consisting of 70,000 grayscale images in 10 categories.
 - b. **Street View House Numbers (SVHN):** Real-world images for digit recognition.
 - c. **Intel Image Classification:** Images categorized into six classes of natural scenes.
 - d. **Tiny ImageNet:** A subset of the ImageNet dataset with 200 classes.
2. **Data Preparation:**
 - a. **Download and Load the Dataset:**
 - i. Follow the dataset link to download the data.
 - ii. Use appropriate PyTorch Dataset classes or create custom ones if necessary.
 - b. **Preprocessing:**
 - i. Normalize the images using mean and standard deviation suitable for the dataset.
 - ii. Resize images if needed to match input dimensions expected by your model.
 - c. **Data Augmentation:**
 - i. Apply transformations such as random horizontal flips, rotations, or color jitter to increase dataset variability.
3. **Model Building:**

a. Architecture:

- i. Construct a CNN model incorporating the key components discussed in the notebook:
 1. Convolutional layers
 2. Pooling layers
 3. Activation functions (e.g., ReLU)
 4. Fully connected layers
 5. Dropout layers to prevent overfitting

b. Model Complexity:

- i. Ensure the model is appropriately complex for the chosen dataset.

4. Training:

a. Hyperparameters:

- i. Set initial hyperparameters like learning rate, batch size, number of epochs, etc.

b. Optimizer and Loss Function:

- i. Use an appropriate optimizer (e.g., Adam, SGD) and loss function (e.g., CrossEntropyLoss).

c. Training Loop:

- i. Implement the training loop with forward pass, loss computation, backward pass, and weight updates.
- ii. Include validation at the end of each epoch to monitor performance.

d. Learning Rate Scheduling:

- i. Implement a learning rate scheduler to adjust the learning rate during training.

5. Evaluation:

a. Test the Model:

- i. Evaluate your model on the test set.
- ii. Calculate accuracy and other relevant metrics (e.g., precision, recall, F1-score).

b. Confusion Matrix:

- i. Plot a confusion matrix to visualize model performance across classes.

6. Visualization:

a. Filters and Feature Maps:

- i. Visualize the learned filters of the convolutional layers.
- ii. Display feature maps after convolutional layers for sample images.

b. Training Curves:

- i. Plot training and validation loss curves over epochs.

7. Hyperparameter Tuning:

a. Experimentation:

- i. Try different hyperparameters to improve model performance.
- ii. Optionally, use GridSearchCV or similar methods for systematic tuning.

b. Document Findings:

- i. Record how changes in hyperparameters affect the model's performance.

8. Report:

a. Summary:

- i. Write a brief report summarizing your approach, experiments, and findings.

b. Analysis:

- i. Discuss the model's performance and any challenges encountered.
- ii. Interpret the visualizations of filters and feature maps.

Datasets:

1. Fashion-MNIST

- **Description:** A dataset consisting of 70,000 grayscale images in 10 categories, each image is 28x28 pixels. The classes include different types of clothing items.
- **Link:** [Fashion-MNIST GitHub](#)

2. Street View House Numbers (SVHN)

- **Description:** A real-world image dataset for developing machine learning and object recognition algorithms with minimal requirement on data preprocessing and formatting. It has over 600,000 digit images (from 0 to 9).
- **Link:** [SVHN Dataset](#)

3. Intel Image Classification

- **Description:** Contains around 25,000 images categorized into six classes: buildings, forest, glacier, mountain, sea, and street.
- **Link:** Intel Image Classification Dataset:
<https://www.kaggle.com/datasets/puneet6060/intel-image-classification>

4. Tiny ImageNet

- **Description:** A subset of the ImageNet dataset with 200 classes, each containing 500 training images, 50 validation images, and 50 test images.
- **Link:** Tiny ImageNet Visual Recognition Challenge
<https://www.kaggle.com/c/tiny-imagenet/data>

Deliverables:

1. **Jupyter Notebook:**
 - a. Well-documented code with explanations and comments.
 - b. Outputs showing results, visualizations, and training progress.
2. **Report Document:**
 - a. A concise report summarizing your methodology, experiments, results, and interpretations.
 - b. Include graphs, charts, and images where relevant.